

## import libraries

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

## Import dataset

```
In [2]: data=pd.read_csv(r"C:\Users\user\Desktop\vicky\C10_air\csvs_per_year\csvs_per_year\madrid_2016.csv")
```

```
In [3]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 500 entries, 0 to 499
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   date        500 non-null    object
 1   BEN         126 non-null    float64
 2   CO          209 non-null    float64
 3   EBE         126 non-null    float64
 4   NMHC        63 non-null     float64
 5   NO          500 non-null    float64
 6   NO_2        500 non-null    float64
 7   O_3         291 non-null    float64
 8   PM10        250 non-null    float64
 9   PM25        126 non-null    float64
10   SO_2        188 non-null    float64
11   TCH         63 non-null     float64
12   TOL         126 non-null    float64
13   station     500 non-null    int64
dtypes: float64(12), int64(1), object(1)
memory usage: 54.8+ KB
```

```
In [4]: data.head()
```

Out[4]:

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL	station
0	2016-11-01 01:00:00	NaN	0.7	NaN	NaN	153.0	77.0	NaN	NaN	NaN	7.0	NaN	NaN	28079004
1	2016-11-01 01:00:00	3.1	1.1	2.0	0.53	260.0	144.0	4.0	46.0	24.0	18.0	2.44	14.4	28079008
2	2016-11-01 01:00:00	5.9	NaN	7.5	NaN	297.0	139.0	NaN	NaN	NaN	NaN	NaN	26.0	28079011
3	2016-11-01 01:00:00	NaN	1.0	NaN	NaN	154.0	113.0	2.0	NaN	NaN	NaN	NaN	NaN	28079016
4	2016-11-01 01:00:00	NaN	NaN	NaN	NaN	275.0	127.0	2.0	NaN	NaN	18.0	NaN	NaN	28079017

```
In [5]: data.shape
```

Out[5]: (500, 14)

```
In [6]: data.index
```

Out[6]: RangeIndex(start=0, stop=500, step=1)

In [7]: data.columns

Out[7]: Index(['date', 'BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO\_2', 'O\_3', 'PM10', 'PM25', 'SO\_2', 'TCH', 'TOL', 'station'], dtype='object')

In [8]: data.isna()

Out[8]:

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL	station
0	False	True	False	True	True	False	False	True	True	True	False	True	True	False
1	False	False	False	False	False	False	False	False	False	False	False	False	False	False
2	False	False	True	False	True	False	False	True	True	True	True	True	False	False
3	False	True	False	True	True	False	False	False	True	True	True	True	True	False
4	False	True	True	True	True	False	False	False	True	True	False	True	True	False
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
495	False	True	True	True	True	False	False	False	True	True	True	True	True	False
496	False	True	True	True	True	False	False	True	False	False	True	True	True	False
497	False	True	True	True	True	False	False	False	True	True	True	True	True	False
498	False	False	True	False	False	False	False	True	False	True	True	False	False	False
499	False	True	False	True	True	False	False	False	True	True	True	True	True	False

500 rows × 14 columns

In [9]: data.fillna(value=0)

Out[9]:

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL	station
0	2016-11-01 01:00:00	0.0	0.7	0.0	0.00	153.0	77.0	0.0	0.0	0.0	7.0	0.00	0.0	28079004
1	2016-11-01 01:00:00	3.1	1.1	2.0	0.53	260.0	144.0	4.0	46.0	24.0	18.0	2.44	14.4	28079008
2	2016-11-01 01:00:00	5.9	0.0	7.5	0.00	297.0	139.0	0.0	0.0	0.0	0.0	0.00	26.0	28079011
3	2016-11-01 01:00:00	0.0	1.0	0.0	0.00	154.0	113.0	2.0	0.0	0.0	0.0	0.00	0.0	28079016
4	2016-11-01 01:00:00	0.0	0.0	0.0	0.00	275.0	127.0	2.0	0.0	0.0	18.0	0.00	0.0	28079017
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
495	2016-11-01 21:00:00	0.0	0.0	0.0	0.00	2.0	64.0	11.0	0.0	0.0	0.0	0.00	0.0	28079049
496	2016-11-01 21:00:00	0.0	0.0	0.0	0.00	22.0	84.0	0.0	28.0	20.0	0.0	0.00	0.0	28079050
497	2016-11-01 21:00:00	0.0	0.0	0.0	0.00	247.0	151.0	3.0	0.0	0.0	0.0	0.00	0.0	28079054
498	2016-11-01 21:00:00	2.2	0.0	1.7	0.30	134.0	106.0	0.0	45.0	0.0	0.0	1.45	8.7	28079055
499	2016-11-01 21:00:00	0.0	1.7	0.0	0.00	278.0	161.0	8.0	0.0	0.0	0.0	0.00	0.0	28079056

500 rows × 14 columns

In [10]: data.isna

Out[10]: <bound method DataFrame.isna of

	date	BEN	CO	EBE	NMHC	NO	NO_2
0	2016-11-01 01:00:00	NaN	0.7	NaN	NaN	153.0	77.0
1	2016-11-01 01:00:00	3.1	1.1	2.0	0.53	260.0	144.0
2	2016-11-01 01:00:00	5.9	NaN	7.5	NaN	297.0	139.0
3	2016-11-01 01:00:00	NaN	1.0	NaN	NaN	154.0	113.0
4	2016-11-01 01:00:00	NaN	NaN	NaN	NaN	275.0	127.0
..	...	...	...	...	...	...	...
495	2016-11-01 21:00:00	NaN	NaN	NaN	NaN	2.0	64.0
496	2016-11-01 21:00:00	NaN	NaN	NaN	NaN	22.0	84.0
497	2016-11-01 21:00:00	NaN	NaN	NaN	NaN	247.0	151.0
498	2016-11-01 21:00:00	2.2	NaN	1.7	0.30	134.0	106.0
499	2016-11-01 21:00:00	NaN	1.7	NaN	NaN	278.0	161.0

  

	SO_2	TCH	TOL	station
0	7.0	NaN	NaN	28079004
1	18.0	2.44	14.4	28079008
2	NaN	NaN	26.0	28079011
3	NaN	NaN	NaN	28079016
4	18.0	NaN	NaN	28079017
..	...	...	...	...
495	NaN	NaN	NaN	28079049
496	NaN	NaN	NaN	28079050
497	NaN	NaN	NaN	28079054
498	NaN	1.45	8.7	28079055
499	NaN	NaN	NaN	28079056

[500 rows x 14 columns]>

## Plotting using various method

In [11]: data.plot.line()

Out[11]: <AxesSubplot:>

```
In [12]: data.plot.bar()
```

```
Out[12]: <AxesSubplot:>
```

```
In [13]: data.plot.area()
```

```
Out[13]: <AxesSubplot:>
```

```
In [14]: data.plot.hist()
```

```
Out[14]: <AxesSubplot:ylabel='Frequency'>
```

```
In [15]: data.plot.pie(y="BEN")
```

```
Out[15]: <AxesSubplot:ylabel='BEN'>
```

```
In [16]: data.plot.scatter(x="NO_2",y='O_3')
```

```
Out[16]: <AxesSubplot:xlabel='NO_2', ylabel='O_3'>
```

## seaborn Visualize

```
In [17]: sns.pairplot(data)
```

```
Out[17]: <seaborn.axisgrid.PairGrid at 0x21089d854f0>
```

```
In [18]: sns.distplot(data['BEN'])
```

```
C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
  warnings.warn(msg, FutureWarning)
```

```
Out[18]: <AxesSubplot:xlabel='BEN', ylabel='Density'>
```

```
In [19]: sns.heatmap(data.corr())
```

```
Out[19]: <AxesSubplot:>
```

```
In [20]: data1=data[['BEN', 'CO', 'EBE', 'NMHC', 'NO_2', 'O_3',  
                  'PM10', 'SO_2']]
```

```
In [21]: data2=data1.fillna(value=1)
```

```
In [22]: x=data2[['CO', 'CO', 'O_3']]  
         y=data['station']
```

## Linear Regression

```
In [23]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [24]: from sklearn.linear_model import LinearRegression
lr=LinearRegression()
lr.fit(x_train,y_train)
```

Out[24]: LinearRegression()

```
In [25]: print(lr.intercept_)
```

28079022.144271646

```
In [26]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['PM10'])
coeff
```

Out[26]:

	PM10
CO	8.817478
CO	8.817478
O_3	0.079522

```
In [27]: prediction1=lr.predict(x_train)
plt.scatter(y_train,prediction1)
```

Out[27]: <matplotlib.collections.PathCollection at 0x21092c44c10>

```
In [28]: lr.score(x_test,y_test)
```

Out[28]: 0.16365927631613542

```
In [29]: prediction1=lr.predict(x_test)
```

## Ridge



```
In [30]: from sklearn.linear_model import Ridge,Lasso  
rr=Ridge(alpha=10)  
rr.fit(x_train,y_train)
```

Out[30]: Ridge(alpha=10)

```
In [31]: rr.score(x_test,y_test)
```

Out[31]: 0.1519329354504101

```
In [32]: prediction2=rr.predict(x_test)
```

## Lasso

```
In [33]: la=Lasso(alpha=10)  
la.fit(x_train,y_train)
```

Out[33]: Lasso(alpha=10)

```
In [34]: la.score(x_test,y_test)
```

Out[34]: -0.0018616480477133823

```
In [35]: prediction3=la.score(x_test,y_test)
```

## Elastic Net

```
In [36]: from sklearn.linear_model import ElasticNet  
en=ElasticNet()  
en.fit(x_train,y_train)
```

Out[36]: ElasticNet()

```
In [37]: print(en.coef_)
```

[1.83024402 1.83024537 0.03948572]

```
In [38]: print(en.intercept_)
```

28079033.73791015

```
In [39]: prediction4=en.predict(x_test)
```

```
In [40]: en.score(x_test,y_test)
```

Out[40]: 0.04607888948793715

## Evaluation Metrics for linear

```
In [41]: from sklearn import metrics
```

```
In [42]: print("Mean Absolute error:",metrics.mean_absolute_error(y_test,prediction1))
```

Mean Absolute error: 14.451300015151501

```
In [43]: print("Mean Absolute square error:",metrics.mean_squared_error(y_test,prediction1))
```

Mean Absolute square error: 277.9946013374757

## Evaluation Metrics for Ridge

```
In [44]: print("Mean Absolute error:",metrics.mean_absolute_error(y_test,prediction2))
```

Mean Absolute error: 14.621402631849051

```
In [45]: print("Mean Absolute square error:",metrics.mean_squared_error(y_test,prediction2))
```

Mean Absolute square error: 281.89236616202686

## Evaluation for elasticnet

```
In [46]: print("Mean Absolute error:",metrics.mean_absolute_error(y_test,prediction4))
```

Mean Absolute error: 15.68005603817602

```
In [47]: print("Mean Absolute square error:",metrics.mean_squared_error(y_test,prediction4))
```

Mean Absolute square error: 317.07761121105284

## Feature matrix

```
In [48]: from sklearn.preprocessing import StandardScaler
```

```
In [49]: from sklearn import utils
```

```
In [50]: from sklearn.linear_model import LogisticRegression
```

```
In [51]: df=pd.read_csv(r"C:\Users\user\Desktop\vicky\C10_air\csvs_per_year\csvs_per_year\madrid_2017.csv")
```

```
In [52]: df.columns
```

```
Out[52]: Index(['date', 'BEN', 'CH4', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'NOx', 'O_3',  
              'PM10', 'PM25', 'SO_2', 'TCH', 'TOL', 'station'],  
              dtype='object')
```

```
In [53]: new_df=df.fillna({'BEN':1,'CO':2,'EBE':4})
new_df
```

Out[53]:

	date	BEN	CH4	CO	EBE	NMHC	NO	NO_2	NOx	O_3	PM10	PM25	SO_2	TCH	TOL	station
0	2017-06-01 01:00:00	1.0	NaN	0.3	4.0	NaN	4.0	38.0	NaN	NaN	NaN	NaN	5.0	NaN	NaN	28079004
1	2017-06-01 01:00:00	0.6	NaN	0.3	0.4	0.08	3.0	39.0	NaN	71.0	22.0	9.0	7.0	1.4	2.9	28079008
2	2017-06-01 01:00:00	0.2	NaN	2.0	0.1	NaN	1.0	14.0	NaN	NaN	NaN	NaN	NaN	NaN	0.9	28079011
3	2017-06-01 01:00:00	1.0	NaN	0.2	4.0	NaN	1.0	9.0	NaN	91.0	NaN	NaN	NaN	NaN	NaN	28079016
4	2017-06-01 01:00:00	1.0	NaN	2.0	4.0	NaN	1.0	19.0	NaN	69.0	NaN	NaN	2.0	NaN	NaN	28079017
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
210115	2017-08-01 00:00:00	1.0	NaN	0.2	4.0	NaN	1.0	27.0	NaN	65.0	NaN	NaN	NaN	NaN	NaN	28079056
210116	2017-08-01 00:00:00	1.0	NaN	0.2	4.0	NaN	1.0	14.0	NaN	NaN	73.0	NaN	7.0	NaN	NaN	28079057
210117	2017-08-01 00:00:00	1.0	NaN	2.0	4.0	NaN	1.0	4.0	NaN	83.0	NaN	NaN	NaN	NaN	NaN	28079058
210118	2017-08-01 00:00:00	1.0	NaN	2.0	4.0	NaN	1.0	11.0	NaN	78.0	NaN	NaN	NaN	NaN	NaN	28079059
210119	2017-08-01 00:00:00	1.0	NaN	2.0	4.0	NaN	1.0	14.0	NaN	77.0	60.0	NaN	NaN	NaN	NaN	28079060

210120 rows × 16 columns

```
In [54]: feature_matrix = new_df[['CO','EBE']]
target_vector = new_df['station']
```

```
In [55]: feature_matrix.shape
```

Out[55]: (210120, 2)

```
In [56]: target_vector.shape
```

Out[56]: (210120,)

```
In [57]: from sklearn.preprocessing import StandardScaler
```

```
In [58]: fs = StandardScaler().fit_transform(feature_matrix)
```

```
In [59]: logr=LogisticRegression()
```

```
In [60]: logr.fit(fs,target_vector)
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:763: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression) ([https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression))

```
n_iter_i = _check_optimize_result(
```

```
Out[60]: LogisticRegression()
```

```
In [61]: observation =[[3,90]]
```

```
In [62]: prediction5 =logr.predict(observation)
print(prediction5)
```

```
[28079016]
```

```
In [63]: logr.predict_proba(observation)[0][0]
```

```
Out[63]: 0.005714299680349634
```

```
In [64]: logr.predict_proba(observation)[0][1]
```

```
Out[64]: 7.251029079565403e-172
```

## import pickle

```
In [65]: import pickle
```

```
In [66]: filename1="prediction1"
```

```
In [67]: filename2="prediction2"
```

```
In [68]: filename3="prediction3"
```

```
In [69]: filename4="prediction4"
```

```
In [70]: filename5="prediction5"
```

```
In [71]: pickle.dump(lr,open(filename1,'wb'))
```

```
In [72]: pickle.dump(lr,open(filename2,'wb'))
```

```
In [73]: pickle.dump(lr,open(filename3,'wb'))
```

```
In [74]: pickle.dump(lr,open(filename4,'wb'))
```

```
In [75]: pickle.dump(lr,open(filename5,'wb'))
```

```
In [ ]:
```