

import libraries

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

Import dataset

```
In [2]: data=pd.read_csv(r"C:\Users\user\Desktop\vicky\C10_air\csvs_per_year\csvs_per_year\madrid_2018.csv")
```

```
In [3]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 500 entries, 0 to 499
Data columns (total 16 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   date        500 non-null    object
 1   BEN         126 non-null    float64
 2   CH4         63 non-null     float64
 3   CO          206 non-null    float64
 4   EBE         126 non-null    float64
 5   NMHC        63 non-null     float64
 6   NO          495 non-null    float64
 7   NO_2        495 non-null    float64
 8   NOx         495 non-null    float64
 9   O_3         286 non-null    float64
10  PM10        271 non-null    float64
11  PM25        147 non-null    float64
12  SO_2        209 non-null    float64
13  TCH         63 non-null     float64
14  TOL         126 non-null    float64
15  station     500 non-null    int64
dtypes: float64(14), int64(1), object(1)
memory usage: 62.6+ KB
```

```
In [4]: data.head()
```

Out[4]:

	date	BEN	CH4	CO	EBE	NMHC	NO	NO_2	NOx	O_3	PM10	PM25	SO_2	TCH	TOL	station
0	2018-03-01 01:00:00	NaN	NaN	0.3	NaN	NaN	1.0	29.0	31.0	NaN	NaN	NaN	2.0	NaN	NaN	28079004
1	2018-03-01 01:00:00	0.5	1.39	0.3	0.2	0.02	6.0	40.0	49.0	52.0	5.0	4.0	3.0	1.41	0.8	28079008
2	2018-03-01 01:00:00	0.4	NaN	NaN	0.2	NaN	4.0	41.0	47.0	NaN	NaN	NaN	NaN	NaN	1.1	28079011
3	2018-03-01 01:00:00	NaN	NaN	0.3	NaN	NaN	1.0	35.0	37.0	54.0	NaN	NaN	NaN	NaN	NaN	28079016
4	2018-03-01 01:00:00	NaN	NaN	NaN	NaN	NaN	1.0	27.0	29.0	49.0	NaN	NaN	3.0	NaN	NaN	28079017

In [5]: data.shape

Out[5]: (500, 16)

In [6]: data.index

Out[6]: RangeIndex(start=0, stop=500, step=1)

In [7]: data.columns

Out[7]: Index(['date', 'BEN', 'CH4', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'NOx', 'O_3',
'PM10', 'PM25', 'SO_2', 'TCH', 'TOL', 'station'],
dtype='object')

In [8]: data.isna()

Out[8]:

	date	BEN	CH4	CO	EBE	NMHC	NO	NO_2	NOx	O_3	PM10	PM25	SO_2	TCH	TOL	station
0	False	True	True	False	True	True	False	False	False	True	True	True	False	True	True	False
1	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False
2	False	False	True	True	False	True	False	False	False	True	True	True	True	True	False	False
3	False	True	True	False	True	True	False	False	False	False	True	True	True	True	True	False
4	False	True	True	True	True	True	False	False	False	False	True	True	False	True	True	False
...
495	False	True	True	True	True	True	False	False	False	False	True	True	True	True	True	False
496	False	True	True	True	True	True	False	False	False	True	False	False	True	True	True	False
497	False	True	True	True	True	True	False	False	False	False	True	True	True	True	True	False
498	False	False	False	True	False	False	False	False	False	True	False	True	True	False	False	False
499	False	True	True	False	True	True	False	False	False	False	False	False	True	True	True	False

500 rows × 16 columns

In [9]: data.fillna(value=0)

Out[9]:

	date	BEN	CH4	CO	EBE	NMHC	NO	NO_2	NOx	O_3	PM10	PM25	SO_2	TCH	TOL	station
0	2018-03-01 01:00:00	0.0	0.00	0.3	0.0	0.00	1.0	29.0	31.0	0.0	0.0	0.0	2.0	0.00	0.0	28079004
1	2018-03-01 01:00:00	0.5	1.39	0.3	0.2	0.02	6.0	40.0	49.0	52.0	5.0	4.0	3.0	1.41	0.8	28079008
2	2018-03-01 01:00:00	0.4	0.00	0.0	0.2	0.00	4.0	41.0	47.0	0.0	0.0	0.0	0.0	0.00	1.1	28079011
3	2018-03-01 01:00:00	0.0	0.00	0.3	0.0	0.00	1.0	35.0	37.0	54.0	0.0	0.0	0.0	0.00	0.0	28079016
4	2018-03-01 01:00:00	0.0	0.00	0.0	0.0	0.00	1.0	27.0	29.0	49.0	0.0	0.0	3.0	0.00	0.0	28079017
...
495	2018-03-01 21:00:00	0.0	0.00	0.0	0.0	0.00	1.0	18.0	19.0	66.0	0.0	0.0	0.0	0.00	0.0	28079049
496	2018-03-01 21:00:00	0.0	0.00	0.0	0.0	0.00	30.0	35.0	81.0	0.0	6.0	5.0	0.0	0.00	0.0	28079050
497	2018-03-01 21:00:00	0.0	0.00	0.0	0.0	0.00	1.0	18.0	20.0	67.0	0.0	0.0	0.0	0.00	0.0	28079054
498	2018-03-01 21:00:00	0.4	1.20	0.0	0.1	0.06	2.0	22.0	26.0	0.0	4.0	0.0	0.0	1.25	0.8	28079055
499	2018-03-01 21:00:00	0.0	0.00	0.3	0.0	0.00	19.0	35.0	64.0	65.0	8.0	4.0	0.0	0.00	0.0	28079056

500 rows × 16 columns

In [10]: data.isna

Out[10]: <bound method DataFrame.isna of

	date	BEN	CH4	CO	EBE	NMHC	NO	NO_2	NOx	O_3	PM10	PM25	SO_2	TCH	TOL	station
0	2018-03-01 01:00:00	0.0	0.00	0.3	0.0	0.00	1.0	29.0	31.0	0.0	0.0	0.0	2.0	0.00	0.0	28079004
1	2018-03-01 01:00:00	0.5	1.39	0.3	0.2	0.02	6.0	40.0	49.0	52.0	5.0	4.0	3.0	1.41	0.8	28079008
2	2018-03-01 01:00:00	0.4	0.00	0.0	0.2	0.00	4.0	41.0	47.0	0.0	0.0	0.0	0.0	0.00	1.1	28079011
3	2018-03-01 01:00:00	0.0	0.00	0.3	0.0	0.00	1.0	35.0	37.0	54.0	0.0	0.0	0.0	0.00	0.0	28079016
4	2018-03-01 01:00:00	0.0	0.00	0.0	0.0	0.00	1.0	27.0	29.0	49.0	0.0	0.0	3.0	0.00	0.0	28079017
...
495	2018-03-01 21:00:00	0.0	0.00	0.0	0.0	0.00	1.0	18.0	19.0	66.0	0.0	0.0	0.0	0.00	0.0	28079049
496	2018-03-01 21:00:00	0.0	0.00	0.0	0.0	0.00	30.0	35.0	81.0	0.0	6.0	5.0	0.0	0.00	0.0	28079050
497	2018-03-01 21:00:00	0.0	0.00	0.0	0.0	0.00	1.0	18.0	20.0	67.0	0.0	0.0	0.0	0.00	0.0	28079054
498	2018-03-01 21:00:00	0.4	1.20	0.0	0.1	0.06	2.0	22.0	26.0	0.0	4.0	0.0	0.0	1.25	0.8	28079055
499	2018-03-01 21:00:00	0.0	0.00	0.3	0.0	0.00	19.0	35.0	64.0	65.0	8.0	4.0	0.0	0.00	0.0	28079056

[500 rows x 16 columns]>

Plotting using various method

In [11]: `data.plot.line()`

Out[11]: `<AxesSubplot:>`

In [12]: `data.plot.bar()`

Out[12]: `<AxesSubplot:>`

```
In [13]: data.plot.area()
```

```
Out[13]: <AxesSubplot:>
```

```
In [14]: data.plot.hist()
```

```
Out[14]: <AxesSubplot:ylabel='Frequency'>
```

```
In [15]: data.plot.pie(y="BEN")
```

```
Out[15]: <AxesSubplot:ylabel='BEN'>
```

```
In [16]: data.plot.scatter(x="NO_2",y='O_3')
```

```
Out[16]: <AxesSubplot:xlabel='NO_2', ylabel='O_3'>
```

seaborn Visualize

```
In [17]: sns.pairplot(data)
```

```
Out[17]: <seaborn.axisgrid.PairGrid at 0x28bcafd9880>
```

```
In [18]: sns.distplot(data['BEN'])
```

```
C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
  warnings.warn(msg, FutureWarning)
```

```
Out[18]: <AxesSubplot:xlabel='BEN', ylabel='Density'>
```

```
In [19]: sns.heatmap(data.corr())
```

```
Out[19]: <AxesSubplot:>
```

```
In [20]: data1=data[['BEN', 'CO', 'EBE', 'NMHC', 'NO_2', 'O_3',  
                  'PM10', 'SO_2']]
```

```
In [21]: data2=data1.fillna(value=1)
```

```
In [22]: x=data2[['CO', 'CO', 'O_3']]  
         y=data['station']
```

Linear Regression


```
In [23]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [24]: from sklearn.linear_model import LinearRegression
lr=LinearRegression()
lr.fit(x_train,y_train)
```

```
Out[24]: LinearRegression()
```

```
In [25]: print(lr.intercept_)
```

```
28079022.527471773
```

```
In [26]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['PM10'])
coeff
```

```
Out[26]:
```

	PM10
CO	10.310919
CO	10.310919
O_3	0.016526

```
In [27]: prediction1=lr.predict(x_train)
plt.scatter(y_train,prediction1)
```

```
Out[27]: <matplotlib.collections.PathCollection at 0x28bd5cd2fa0>
```

```
In [28]: lr.score(x_test,y_test)
```

```
Out[28]: 0.12718946102785744
```

```
In [29]: prediction1=lr.predict(x_test)
```

Ridge

```
In [30]: from sklearn.linear_model import Ridge,Lasso  
rr=Ridge(alpha=10)  
rr.fit(x_train,y_train)
```

Out[30]: Ridge(alpha=10)

```
In [31]: rr.score(x_test,y_test)
```

Out[31]: 0.1282399636175272

```
In [32]: prediction2=rr.predict(x_test)
```

Lasso

```
In [33]: la=Lasso(alpha=10)  
la.fit(x_train,y_train)
```

Out[33]: Lasso(alpha=10)

```
In [34]: la.score(x_test,y_test)
```

Out[34]: -0.004789289789998374

```
In [35]: prediction3=la.score(x_test,y_test)
```

Elastic Net

```
In [36]: from sklearn.linear_model import ElasticNet  
en=ElasticNet()  
en.fit(x_train,y_train)
```

Out[36]: ElasticNet()

```
In [37]: print(en.coef_)
```

[2.74144059 2.74145716 -0.01452483]

```
In [38]: print(en.intercept_)
```

28079034.231329445

```
In [39]: prediction4=en.predict(x_test)
```

```
In [40]: en.score(x_test,y_test)
```

Out[40]: 0.05936449184919146

Evaluation Metrics for linear

```
In [41]: from sklearn import metrics
```

```
In [42]: print("Mean Absolute error:",metrics.mean_absolute_error(y_test,prediction1))
```

Mean Absolute error: 14.018862525075674

```
In [43]: print("Mean Absolute square error:",metrics.mean_squared_error(y_test,prediction1))
```

Mean Absolute square error: 283.53153379561275

Evaluation Metrics for Ridge

```
In [44]: print("Mean Absolute error:",metrics.mean_absolute_error(y_test,prediction2))
```

Mean Absolute error: 14.112414263164004

```
In [45]: print("Mean Absolute square error:",metrics.mean_squared_error(y_test,prediction2))
```

Mean Absolute square error: 283.1902791965836

Evaluation for elasticnet

```
In [46]: print("Mean Absolute error:",metrics.mean_absolute_error(y_test,prediction4))
```

Mean Absolute error: 14.796460316677889

```
In [47]: print("Mean Absolute square error:",metrics.mean_squared_error(y_test,prediction4))
```

Mean Absolute square error: 305.56439967222553

Feature matrix

```
In [48]: from sklearn.preprocessing import StandardScaler
```

```
In [49]: from sklearn import utils
```

```
In [50]: from sklearn.linear_model import LogisticRegression
```

```
In [51]: df=pd.read_csv(r"C:\Users\user\Desktop\vicky\C10_air\csvs_per_year\csvs_per_year\madrid_2018.csv")
```

```
In [52]: df.columns
```


```
Out[52]: Index(['date', 'BEN', 'CH4', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'NOx', 'O_3',  
              'PM10', 'PM25', 'SO_2', 'TCH', 'TOL', 'station'],  
              dtype='object')
```

```
In [53]: new_df=df.fillna({'BEN':1,'CO':2,'EBE':4})
new_df
```

Out[53]:

	date	BEN	CH4	CO	EBE	NMHC	NO	NO_2	NOx	O_3	PM10	PM25	SO_2	TCH	TOL	station
0	2018-03-01 01:00:00	1.0	NaN	0.3	4.0	NaN	1.0	29.0	31.0	NaN	NaN	NaN	2.0	NaN	NaN	28079004
1	2018-03-01 01:00:00	0.5	1.39	0.3	0.2	0.02	6.0	40.0	49.0	52.0	5.0	4.0	3.0	1.41	0.8	28079008
2	2018-03-01 01:00:00	0.4	NaN	2.0	0.2	NaN	4.0	41.0	47.0	NaN	NaN	NaN	NaN	NaN	1.1	28079011
3	2018-03-01 01:00:00	1.0	NaN	0.3	4.0	NaN	1.0	35.0	37.0	54.0	NaN	NaN	NaN	NaN	NaN	28079016
4	2018-03-01 01:00:00	1.0	NaN	2.0	4.0	NaN	1.0	27.0	29.0	49.0	NaN	NaN	3.0	NaN	NaN	28079017
...
69091	2018-02-01 00:00:00	1.0	NaN	0.5	4.0	NaN	66.0	91.0	192.0	1.0	35.0	22.0	NaN	NaN	NaN	28079056
69092	2018-02-01 00:00:00	1.0	NaN	0.7	4.0	NaN	87.0	107.0	241.0	NaN	29.0	NaN	15.0	NaN	NaN	28079057
69093	2018-02-01 00:00:00	1.0	NaN	2.0	4.0	NaN	28.0	48.0	91.0	2.0	NaN	NaN	NaN	NaN	NaN	28079058
69094	2018-02-01 00:00:00	1.0	NaN	2.0	4.0	NaN	141.0	103.0	320.0	2.0	NaN	NaN	NaN	NaN	NaN	28079059
69095	2018-02-01 00:00:00	1.0	NaN	2.0	4.0	NaN	69.0	96.0	202.0	3.0	26.0	NaN	NaN	NaN	NaN	28079060

69096 rows × 16 columns



```
In [54]: feature_matrix = new_df[['CO','EBE']]
target_vector = new_df['station']
```

```
In [55]: feature_matrix.shape
```

Out[55]: (69096, 2)

```
In [56]: target_vector.shape
```

Out[56]: (69096,)

```
In [57]: from sklearn.preprocessing import StandardScaler
```

```
In [58]: fs = StandardScaler().fit_transform(feature_matrix)
```

```
In [59]: logr=LogisticRegression()
```

```
In [60]: logr.fit(fs,target_vector)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:763: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```
Out[60]: LogisticRegression()
```

```
In [61]: observation =[[3,90]]
```

```
In [62]: prediction5 =logr.predict(observation)  
print(prediction5)
```

```
[28079004]
```

```
In [63]: logr.predict_proba(observation)[0][0]
```

```
Out[63]: 0.9389282324945811
```

```
In [64]: logr.predict_proba(observation)[0][1]
```

```
Out[64]: 4.486888459491602e-203
```

import pickle

```
In [65]: import pickle
```

```
In [66]: filename1="prediction1"
```

```
In [67]: filename2="prediction2"
```

```
In [68]: filename3="prediction3"
```

```
In [69]: filename4="prediction4"
```

```
In [70]: filename5="prediction5"
```

```
In [71]: pickle.dump(lr,open(filename1,'wb'))
```

```
In [72]: pickle.dump(lr,open(filename2,'wb'))
```

```
In [73]: pickle.dump(lr,open(filename3,'wb'))
```

```
In [74]: pickle.dump(lr,open(filename4,'wb'))
```

```
In [75]: pickle.dump(lr,open(filename5,'wb'))
```

```
In [ ]:
```